

```

> #####Assignment 2 Part 2 Falkenberg

> ###File Download

> #IRIS

> iris <- read.csv("C:/Users/chaos/Downloads/iris.csv")

> iris <- iris[,-1]

> # abalone dataset from UCI repository

> # reading the dataset from UCI repository URL

> abalone <- read.csv(url("https://archive.ics.uci.edu/ml/machine-learning-
databases/abalone/abalone.data"), header = FALSE, sep = ";")

> # rename columns

> colnames(abalone) <- c("sex", "length", 'diameter', 'height', 'whole_weight',
'shucked_wieght', 'viscera_wieght', 'shell_weight',
+           'rings' )

> # add new column abalone$age.group with 3 values based on the number of rings

> abalone$age.group <- cut(abalone$rings, br=c(0,8,11,35), labels = c("young", 'adult',
'old'))

> # drop the sex column (categorical variable)

> abalone.norm <- abalone[,-1]

> # optionally normalize

> #normalize <- function(x) {return ((x - min(x)) / (max(x) - min(x))) }

> #abalone.norm[1:7] <- as.data.frame(lapply(abalone.norm[1:7], normalize))

> ##Exercise 1: Naïve Bayes

> ## Call the NaiveBayes Classifier Package e1071, which auto calls the Class package ##

> #IRIS dataset.

> library("e1071")

> #Train classifier

> classifier<-naiveBayes(iris[,1:4], iris[,5])

```

```

> #evaluate classification
> table(predict(classifier, iris[,-5]), iris[,5], dnn=list('predicted','actual'))

      actual
predicted setosa versicolor virginica
setosa     50      0      0
versicolor  0     47      3
virginica   0      3     47

> #examine class means and standard deviations for petal length
> classifier$tables$Petal.Length

      Petal.Length
iris[, 5]  [,1]   [,2]
setosa    1.462 0.1736640
versicolor 4.260 0.4699110
virginica  5.552 0.5518947

> #plot normal distributions at the means of the classes
> #one class
> plot(function(x) dnorm(x, 1.462, 0.1736640), 0, 8, col="red", main="Petal length
distribution for the 3 different species")

> #another class
> curve(dnorm(x, 4.260, 0.4699110), add=TRUE, col="blue")

> #final class
> curve(dnorm(x, 5.552, 0.5518947 ), add=TRUE, col = "green")

> #Abalone dataset 1: height.
> library("e1071")

> classifier<-naiveBayes(abalone.norm[,1:7], abalone.norm[,9])

> table(predict(classifier, abalone.norm[,1:7]), abalone.norm[,9],
dnn=list('predicted','actual'))

```

```

      actual
predicted young adult old
  young 1122 423 145
  adult 266 910 437
  old   19 477 378
> classifier$tables$height

      height
abalone.norm[, 9]  [,1]  [,2]
  young 0.1065956 0.04183039
  adult 0.1516906 0.02984784
  old   0.1648125 0.02935998

> #Take information from line 40; first should be mean then standard deviation

> plot(function(x) dnorm(x, 0.1065956 , 0.04183039), 0, 0.4, ylim = c(0,14), col="red",
main="Height for the 3 different ages")

> curve(dnorm(x, 0.1516906, 0.02984784), add=TRUE, col="blue")

> curve(dnorm(x, 0.1648125, 0.02935998), add=TRUE, col = "green")

> #Abalone dataset subset 2: length.

> library("e1071")

> classifier<-naiveBayes(abalone.norm[,1:7], abalone.norm[,9])

> table(predict(classifier, abalone.norm[,1:7]), abalone.norm[,9],
dnn=list('predicted','actual'))

      actual
predicted young adult old
  young 1122 423 145
  adult 266 910 437
  old   19 477 378

> classifier$tables$length

```

```

length
abalone.norm[, 9]  [,1]  [,2]

young 0.4209915 0.11137474
adult 0.5707182 0.08740980
old 0.5868542 0.08100644

> plot(function(x) dnorm(x, 0.4209915, 0.11137474), 0, 1, ylim = c(0,5), col="red",
main="Length for the 3 different ages")

> curve(dnorm(x, 0.5707182, 0.08740980), add=TRUE, col="blue")

> curve(dnorm(x, 0.5868542, 0.08100644), add=TRUE, col = "green")

> #Abalone dataset subset 3:whole_weight.

> library("e1071")

> classifier<-naiveBayes(abalone.norm[,1:7], abalone.norm[,9])

> table(predict(classifier, abalone.norm[,1:7]), abalone.norm[,9],
dnn=list('predicted','actual'))

actual

predicted young adult old

young 1122 423 145
adult 266 910 437
old 19 477 378

> classifier$tables$whole_weight

whole_weight

abalone.norm[, 9]  [,1]  [,2]

young 0.4323742 0.3060074
adult 0.9850878 0.4264315
old 1.1148922 0.4563715

> plot(function(x) dnorm(x, 0.4323742, 0.3060074), 0,2.5, ylim = c(0,1.5), col="red",
main="Whole_Weight for the 3 different ages")

```

```

> curve(dnorm(x, 0.9850878 , 0.4264315), add=TRUE, col="blue")
> curve(dnorm(x, 1.1148922 , 0.4563715), add=TRUE, col = "green")
> ##Exercise 2: K-Nearest Neighbors (kNN)
> #Abalone Dataset
> # sample 2924 from 4177 (~70%)
> s_abalone <- sample(4177,2924)
> #Abalone.norm.train <- abalone.norm[s_abalone,]
> #abalone.norm.test <- abalone.norm[-s_abalone,]
> ## create train & test sets based on sampled indexes
> abalone.train <- abalone[s_abalone,]
> abalone.test <- abalone[-s_abalone,]
> sqrt(2924)
[1] 54.07402
> k = 55
> # k = 80
> # train model & predict
> library(class)
> KNNpred <- knn(train = abalone.train[2:7], test = abalone.test[2:7], cl =
abalone.train$age.group, k = k)
> # create contingency table/ confusion matrix
> contingency.table <- table(KNNpred, abalone.test$age.group)
> contingency.matrix = as.matrix(contingency.table)
> sum(diag(contingency.matrix))/length(abalone.test$age.group)
[1] 0.6823623
> accuracy <- c()
> ks <- c(35,45,55,65,75,85,95,105)

```

```

> for (k in ks) {
+   KNNpred <- knn(train = abalone.train[2:7], test = abalone.test[2:7], cl =
abalone.train$age.group, k = k)
+   cm = as.matrix(table(Actual=KNNpred, Predicted = abalone.test$age.group,
dnn=list('predicted','actual')))
+   accuracy <- c(accuracy,sum(diag(cm))/length(abalone.test$age.group))
+ }
> plot(ks,accuracy,type = "b", ylim = c(0.66,0.69))
> plot(ks,accuracy,type = "b", ylim = c(0.66,0.75))
> plot(ks,accuracy,type = "b", ylim = c(0.66,0.70))
> #Iris Subset 1.
> s_iris <- sample(150, 105) # (~70% training data)
> iris.train <- iris[s_iris, ]
> iris.test <- iris[-s_iris, ]
> # Verify the sample size
> sqrt(105)
[1] 10.24695
> k = 10
> # Load required library
> library(class)
> # Implement KNN classification
> KNNpred <- knn(train = iris.train[, 1:4], test = iris.test[, 1:4], cl = iris.train$Species, k = k)
> # Create a contingency table (confusion matrix)
> contingency.table <- table(KNNpred, iris.test$Species)
> # Accuracy calculation
> contingency.matrix <- as.matrix(contingency.table)
> accuracy_single <- sum(diag(contingency.matrix)) / length(iris.test$Species)

```

```

> # Initialize accuracy vector and try different k-values

> accuracy <- c()

> ks <- c(9, 11, 13, 15, 17, 19)

> for(k in ks) {

+   KNNpred <- knn(train = iris.train[, 1:4], test = iris.test[, 1:4], cl = iris.train$Species, k = k)

+   cm <- as.matrix(table(Actual = iris.test$Species, Predicted = KNNpred))

+   accuracy <- c(accuracy, sum(diag(cm)) / length(iris.test$Species))

+ }

> plot(ks, accuracy, type = "b", ylim = c(0.93, 0.96), main = "Accuracy vs. k-values", xlab =
"k", ylab = "Accuracy")

> plot(ks, accuracy, type = "b", ylim = c(0.80, 0.96), main = "Accuracy vs. k-values", xlab =
"k", ylab = "Accuracy")

> plot(ks, accuracy, type = "b", ylim = c(0.84, 0.96), main = "Accuracy vs. k-values", xlab =
"k", ylab = "Accuracy")

> plot(ks, accuracy, type = "b", ylim = c(0.86, 0.96), main = "Accuracy vs. k-values", xlab =
"k", ylab = "Accuracy")

> #Iris Subset 2.

> s_iris <- sample(150, 105) # (~70% training data)

> iris.train <- iris[s_iris, ]

> iris.test <- iris[-s_iris, ]

> # Verify the sample size

> sqrt(105)

[1] 10.24695

> k = 10

> # Load required library

> library(class)

> # Implement KNN classification

```

```

> KNNpred <- knn(train = iris.train[, 1:4], test = iris.test[, 1:4], cl = iris.train$Species, k = k)
> # Create a contingency table (confusion matrix)
> contingency.table <- table(KNNpred, iris.test$Species)
> # Accuracy calculation
> contingency.matrix <- as.matrix(contingency.table)
> accuracy_single <- sum(diag(contingency.matrix)) / length(iris.test$Species)
> print(paste("Accuracy for k =", k, ":", accuracy_single))
[1] "Accuracy for k = 10 : 0.955555555555556"
> # Initialize accuracy vector and try different k-values
> accuracy <- c()
> ks <- c(7, 21, 33, 45, 57, 69, 71, 83, 95, 105)
> for(k in ks) {
+   KNNpred <- knn(train = iris.train[, 1:4], test = iris.test[, 1:4], cl = iris.train$Species, k = k)
+   cm <- as.matrix(table(Actual = iris.test$Species, Predicted = KNNpred))
+   accuracy <- c(accuracy, sum(diag(cm)) / length(iris.test$Species))
+ }
> plot(ks, accuracy, type = "b", ylim = c(0.2, 1), main = "Accuracy vs. k-values", xlab = "k",
ylab = "Accuracy")
> ##Exercise 3: K-Means
> #Iris Dataset
> # Plot iris petal length vs. petal width, color by species
> ggplot(iris, aes(x = Petal.Length, y = Petal.Width, colour = Species)) +
+   geom_point()
Error in ggplot(iris, aes(x = Petal.Length, y = Petal.Width, colour = Species)) :
  could not find function "ggplot"
> library(ggplot2)

```



```

> ##Exercise 3: K-Means

> #Iris Dataset

> # Plot iris petal length vs. petal width, color by species

> ggplot(iris, aes(x = Petal.Length, y = Petal.Width, colour = Species)) +
+   geom_point()

> # set seed for random number generator

> set.seed(123)

> # run k-means

> iris.km <- kmeans(iris[, -5], centers = 3)

> assigned.clusters <- as.factor(iris.km$cluster)

> ggplot(iris, aes(x = Petal.Length, y = Petal.Width, colour = assigned.clusters)) +
+   geom_point()

> wss <- c()

> ks <- c(2,3,4,5)

> for (k in ks) {
+   iris.km <- kmeans(iris[, -5], centers = k)
+   wss <- c(wss, iris.km$tot.withinss)
+ }

> plot(ks, wss, type = "b")

> labeled.clusters <- as.character(assigned.clusters)

> labeled.clusters[labeled.clusters == 1] <- "setosa"

> labeled.clusters[labeled.clusters == 2] <- "versicolor"

> labeled.clusters[labeled.clusters == 3] <- "virginica"

> table(labeled.clusters, iris[, 5])

```

labeled.clusters setosa versicolor virginica

setosa	50	0	0
versivolor	0	48	14
virginica	0	2	36

```
> #Iris with different values of k

> # Plot iris petal length vs. petal width, color by species

> ggplot(iris, aes(x = Petal.Length, y = Petal.Width, colour = Species)) +
+   geom_point()

> # set seed for random number generator

> set.seed(234)

> # run k-means

> iris.km <- kmeans(iris[, -5], centers = 3)

> assigned.clusters <- as.factor(iris.km$cluster)

> ggplot(iris, aes(x = Petal.Length, y = Petal.Width, colour = assigned.clusters)) +
+   geom_point()

> wss <- c()

> ks <- c(1,3,5,7,9)

> for (k in ks) {
+   iris.km <- kmeans(iris[, -5], centers = k)
+   wss <- c(wss, iris.km$tot.withinss)
+ }

> plot(ks, wss, type = "b")

> labeled.clusters <- as.character(assigned.clusters)

> labeled.clusters[labeled.clusters==1] <- "setosa"

> labeled.clusters[labeled.clusters==2] <- "versivolor"

> labeled.clusters[labeled.clusters==3] <- "virginica"

> table(labeled.clusters, iris[, 5])
```

labeled.clusters setosa versicolor virginica

setosa	17	4	0
versicolor	33	0	0
virginica	0	46	50

```
> #Abalone Dataset
```

```
> # Plot abalone height vs. whole_weight, colored by age.group
```

```
> ggplot(abalone.norm, aes(x = height, y = whole_weight, colour = age.group)) +  
+ geom_point() +  
+ ggtitle("Height vs. Whole Weight by Age Group")
```

```
> # Set seed
```

```
> set.seed(345)
```

```
> # Run K-means clustering
```

```
> abalone.km <- kmeans(abalone.norm[,-c(9)], centers = 3)
```

```
> # Assign clusters
```

```
> assigned.clusters <- as.factor(abalone.km$cluster)
```

```
> # Plot height vs whole_weight colored by the clusters assigned by K-means
```

```
> ggplot(abalone.norm, aes(x = height, y = whole_weight, colour = assigned.clusters)) +  
+ geom_point() +  
+ ggtitle("Height vs. Whole Weight by K-means Clusters")
```

```
> # Try different values of k to calculate total within-cluster sum of squares (wss)
```

```
> wss <- c() # Within Sum of Squares
```

```
> ks <- c(2, 3, 4, 5)
```

```
> for (k in ks) {
```

```
+ abalone.km <- kmeans(abalone.norm[,-c(9)], centers = k)
```

```
+ wss <- c(wss, abalone.km$tot.withinss) # Add wss for each k
```

```

+ }

> # Plot the total within-cluster sum of squares (Elbow plot)

> plot(ks, wss, type = "b", xlab = "Number of clusters (k)", ylab = "Total within-cluster sum of
squares",

+   main = "Elbow Method for Optimal k")

> # Map cluster numbers to age groups (young, adult, old)

> labeled.clusters <- as.character(assigned.clusters)

> labeled.clusters[labeled.clusters == "1"] <- "young"

> labeled.clusters[labeled.clusters == "2"] <- "adult"

> labeled.clusters[labeled.clusters == "3"] <- "old"

> table(labeled.clusters, abalone.norm[,9])

```

```
labeled.clusters young adult old
```

```
adult 1407 0 0
```

```
old 0 0 693
```

```
young 0 1810 267
```

```

> #Abalone with different values of k

> # Plot abalone diameter vs. length, colored by age.group

> ggplot(abalone.norm, aes(x = diameter, y = length, colour = age.group)) +

+   geom_point() +

+   ggtitle("Diameter vs. Length by Age Group")

> # Set seed

> set.seed(345)

> # Run K-means clustering

> abalone.km <- kmeans(abalone.norm[, -c(9)], centers = 3)

> # Assign clusters

```

```

> assigned.clusters <- as.factor(abalone.km$cluster)

> # Plot length vs diameter colored by the clusters assigned by K-means

> ggplot(abalone.norm, aes(x = diameter, y = length, colour = assigned.clusters)) +
+   geom_point() +
+   ggtitle("Diameter vs. Length by K-means Clusters")

> # Try different values of k to calculate total within-cluster sum of squares (wss)

> wss <- c() # Within Sum of Squares

> ks <- c(1, 3, 5, 7, 9)

> for (k in ks) {
+   abalone.km <- kmeans(abalone.norm[, -c(9)], centers = k)
+   wss <- c(wss, abalone.km$tot.withinss) # Add wss for each k
+ }

> # Plot the total within-cluster sum of squares (Elbow plot)

> plot(ks, wss, type = "b", xlab = "Number of clusters (k)", ylab = "Total within-cluster sum of
squares",
+   main = "Elbow Method for Optimal k")

> # Map cluster numbers to age groups (young, adult, old)

> labeled.clusters <- as.character(assigned.clusters)

> labeled.clusters[labeled.clusters == "1"] <- "young"

> labeled.clusters[labeled.clusters == "2"] <- "adult"

> labeled.clusters[labeled.clusters == "3"] <- "old"

> table(labeled.clusters, abalone.norm[,9])

```

```
labeled.clusters young adult old
```

```
adult 1407 0 0
```

```
old 0 0 693
```

young 0 1810 267

>